

Problem Solving by Searching

Asmaa Elbadrawy
PhD, Lecturer
IFT Program, ASU

A **goal-based agent** acts to reach a specific goal.

If reaching the goal requires
a sequence of actions,
the agent needs to find
a path to the goal. This is
called a **problem-solving agent**.

Example: Path Finding

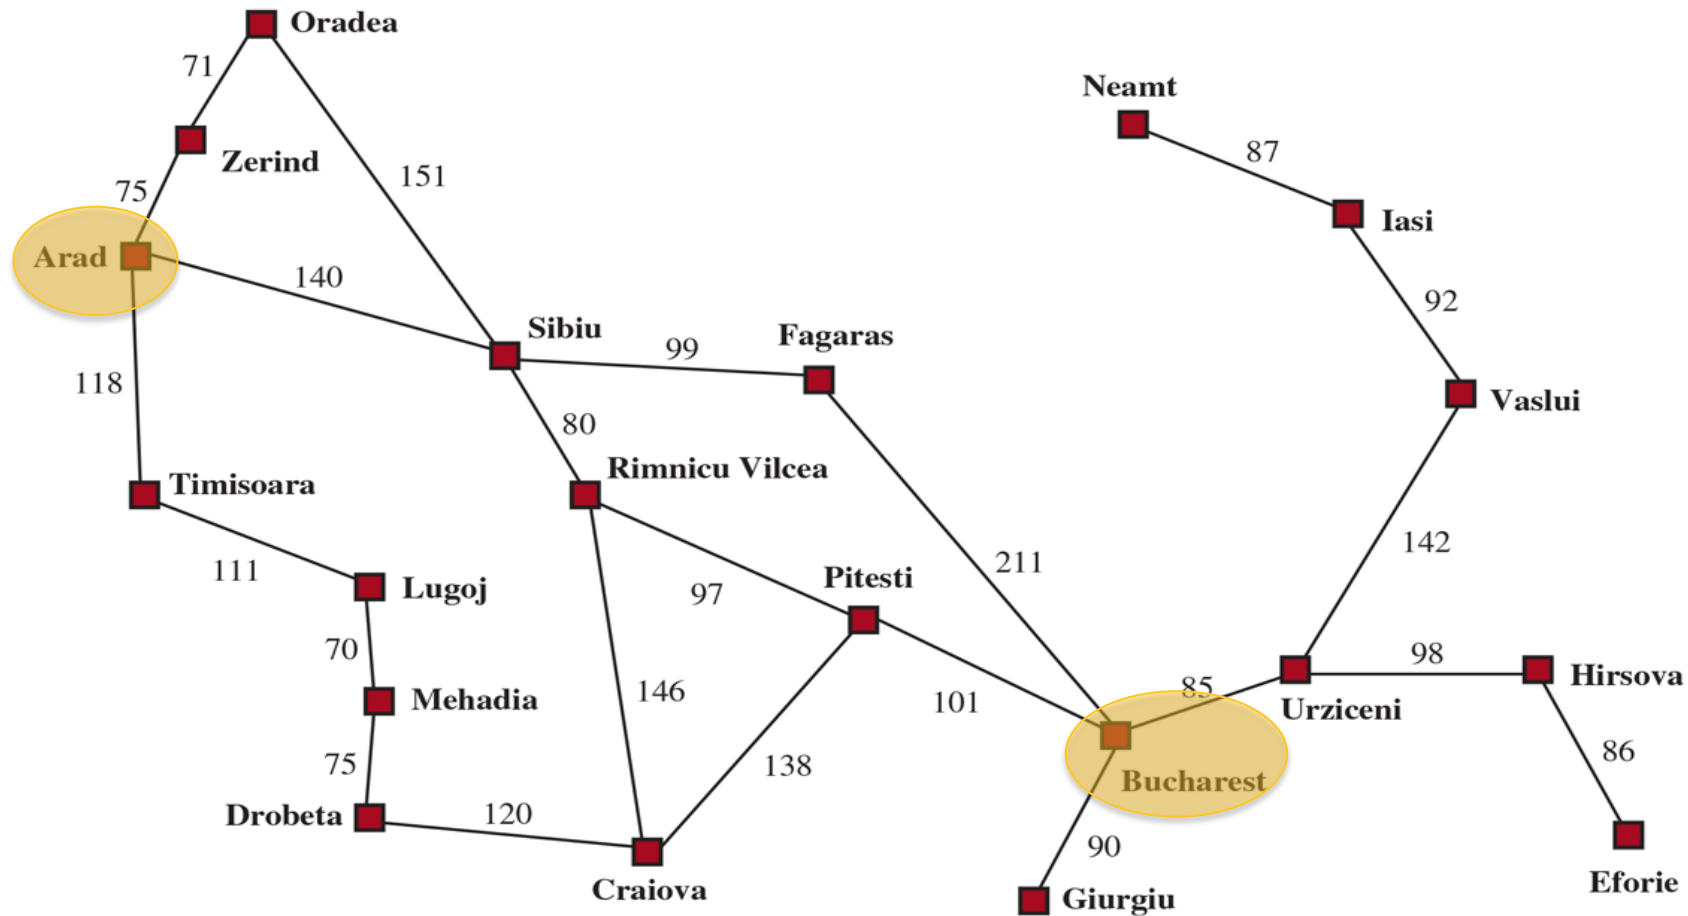


Fig 3.1, Russell & Norvig's Textbook

**Find a path
from Arad to
Bucharest.**

**Similar to
finding
flights with
multiple
legs, or path
finding by
GPS Apps.**

Example: Path Finding


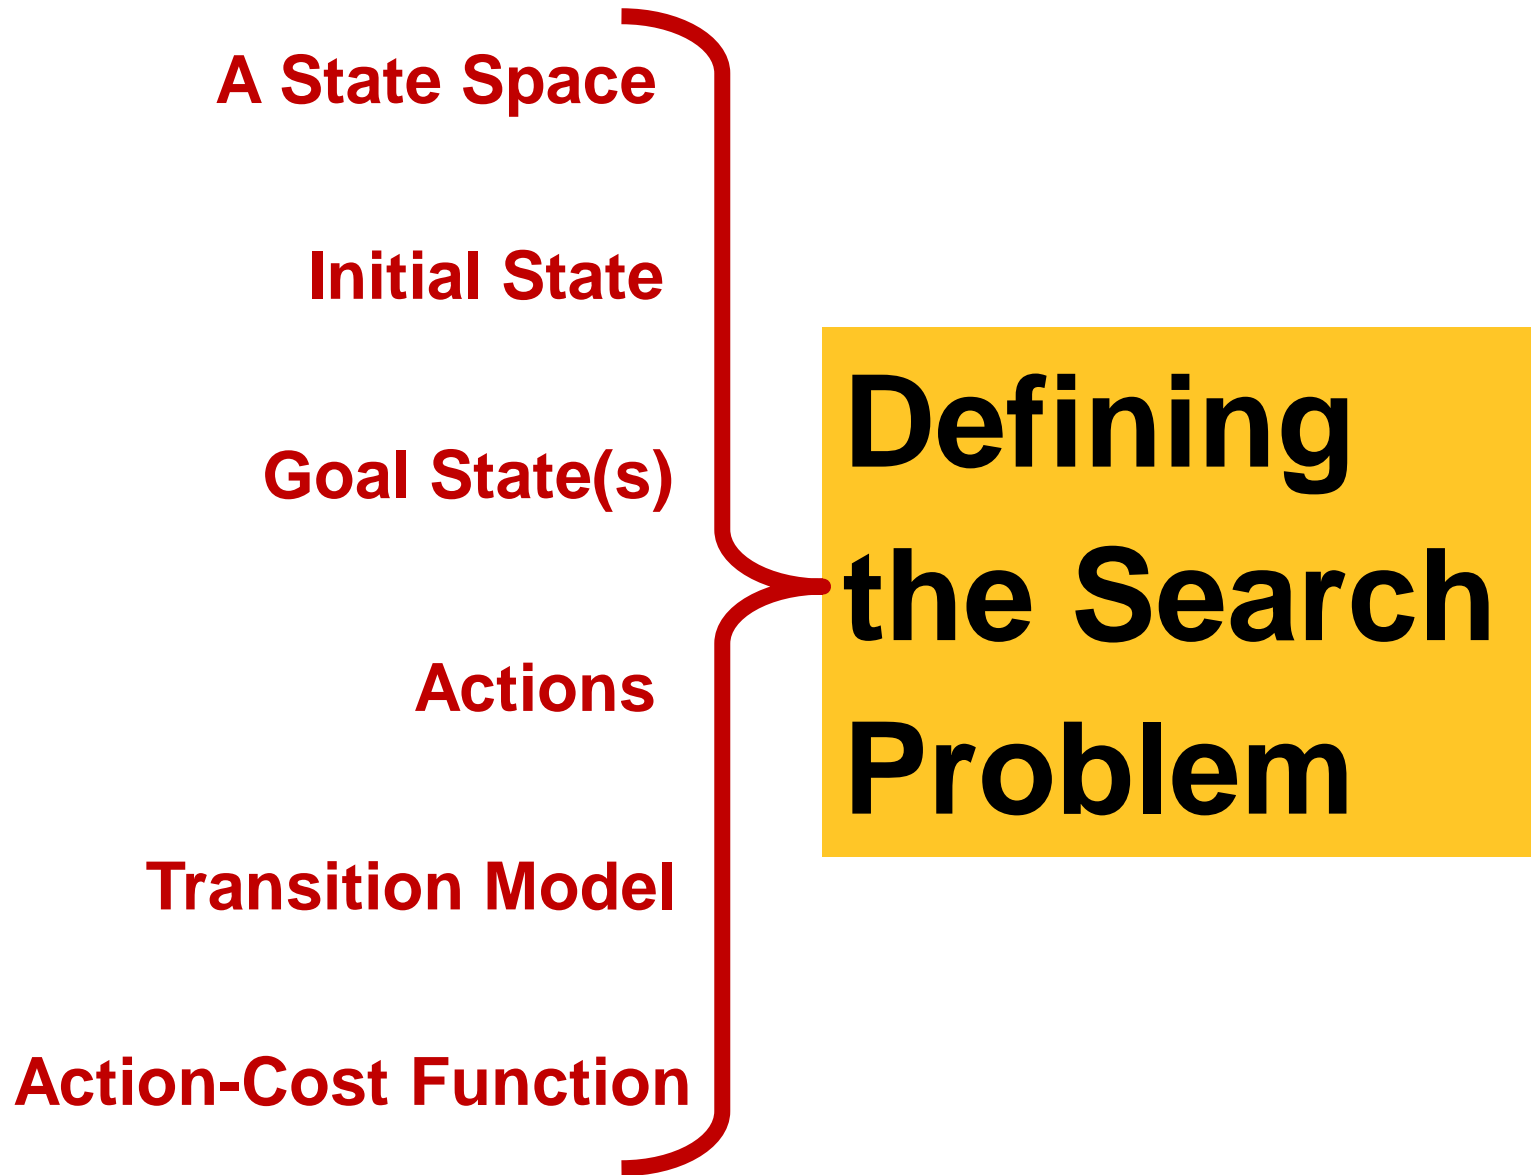


Fig 3.1, Russell & Norvig's Textbook

Since the environment is fully observable, deterministic, and known, the solution is a fixed sequence of actions.

A problem in which we try to find a sequence of actions is called a **search problem**.



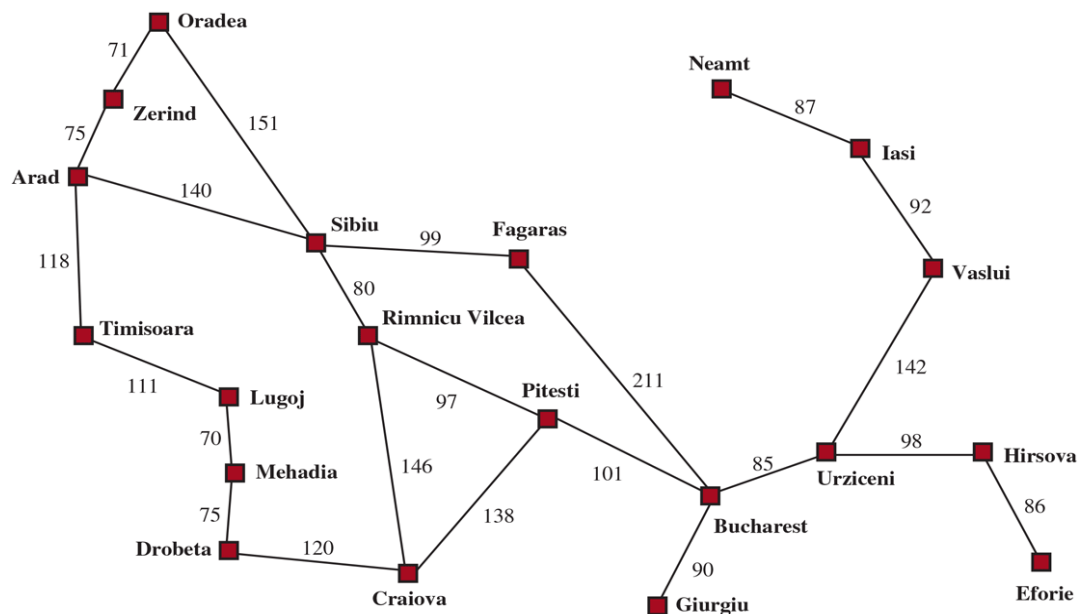


Fig 3.1, Russell & Norvig's Textbook

A State Space

Initial State

Goal State(s)

Actions

Transition Model

Action-Cost Function

**Defining
the Search
Problem**

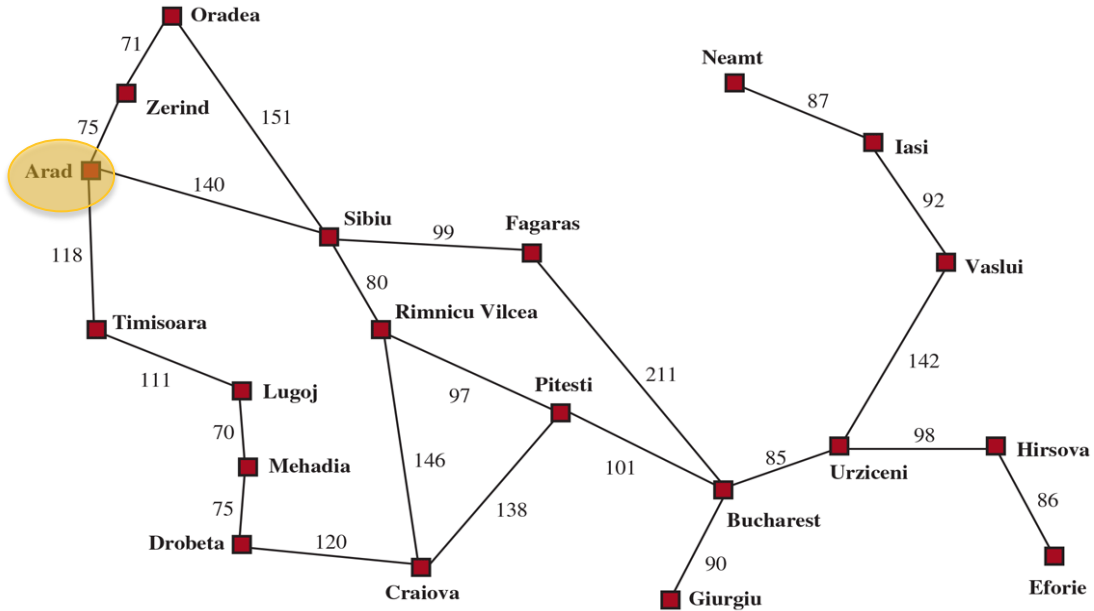


Fig 3.1, Russell & Norvig's Textbook

A State Space

Initial State

Goal State(s)

Actions

Transition Model

Action-Cost Function

Defining
the Search
Problem

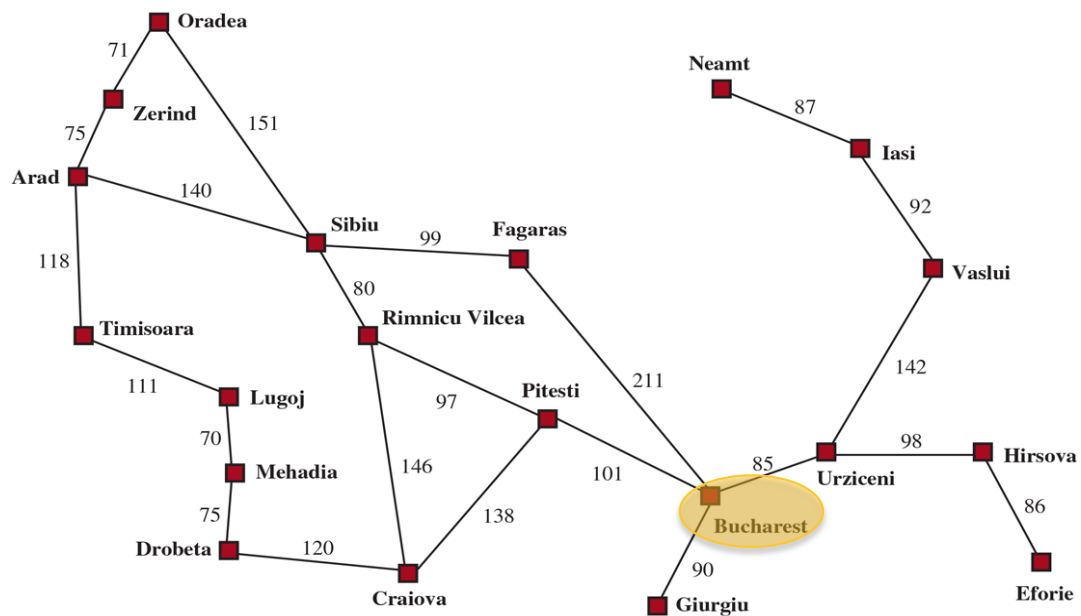


Fig 3.1, Russell & Norvig's Textbook

A State Space

Initial State

Goal State(s)

Actions

Transition Model

Action-Cost Function

Defining
the Search
Problem




Fig 3.1, Russell & Norvig's Textbook

$s = \text{Oradea}$
 $\text{ACTIONS}(s) = \{\text{Zerind}, \text{Sibiu}\}$

A State Space

Initial State

Goal State(s)

Actions

Transition Model

Action-Cost Function

Defining
the Search
Problem

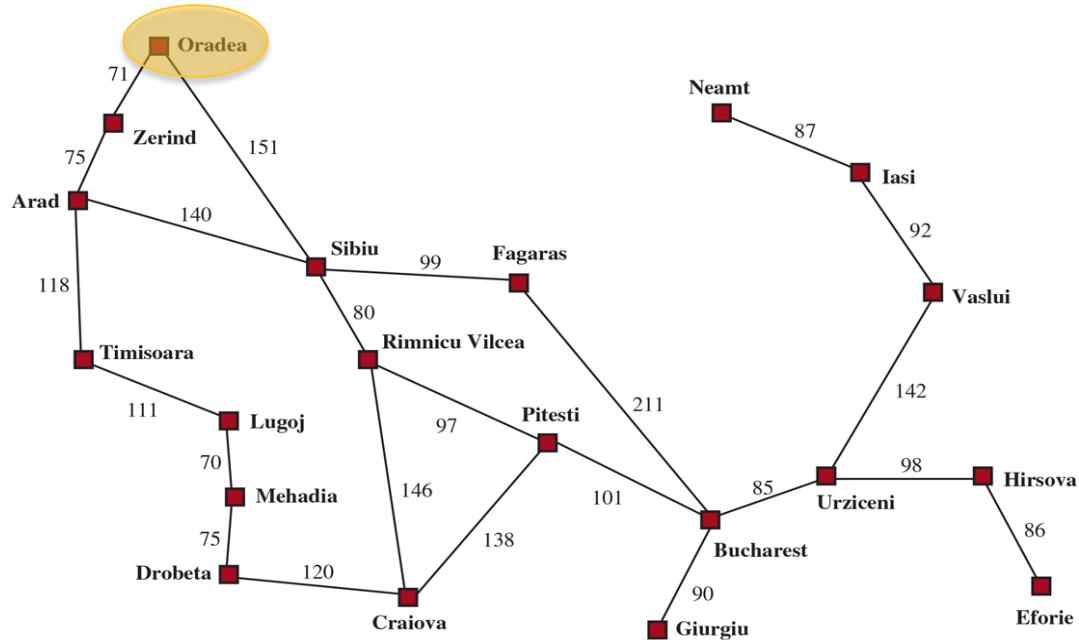


Fig 3.1, Russell & Norvig's Textbook

$s = \text{Oradea}$
 $\text{ACTIONS}(s) = \{\text{ToZerind}, \text{ToSibiu}\}$

A State Space

Initial State

Goal State(s)

Actions

Transition Model

Action-Cost Function

Defining
the Search
Problem

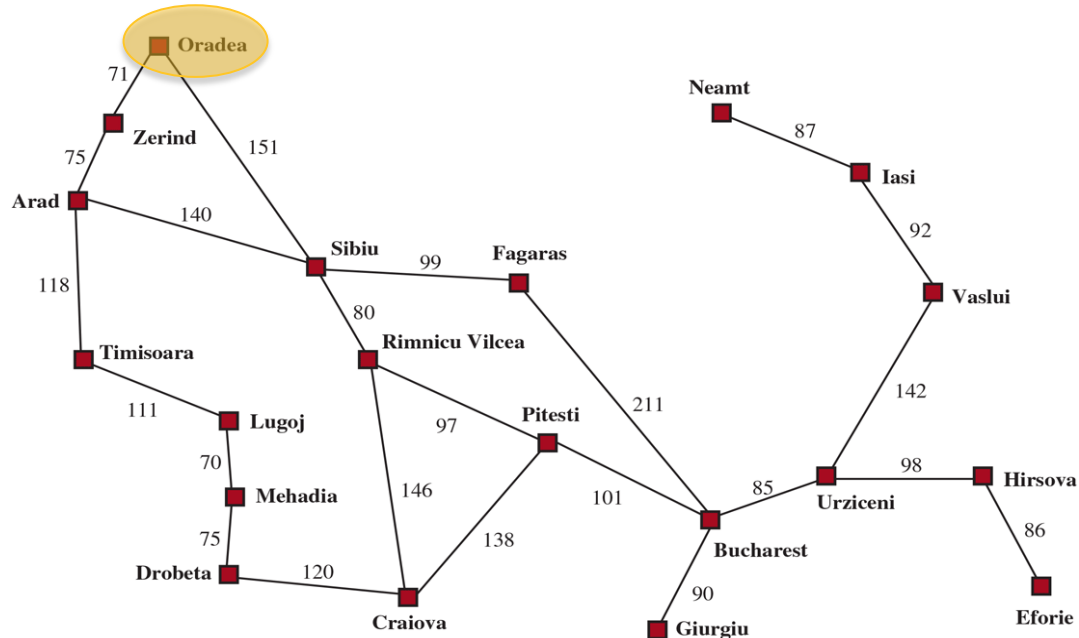


Fig 3.1, Russell & Norvig's Textbook

`s = Oradea, a = ToZerind`
`Result(s,a) = {Oradea,ToZerind}`
`= Zerind`

A State Space

Initial State

Goal State(s)

Actions

Transition Model

Action-Cost Function

Defining
the Search
Problem

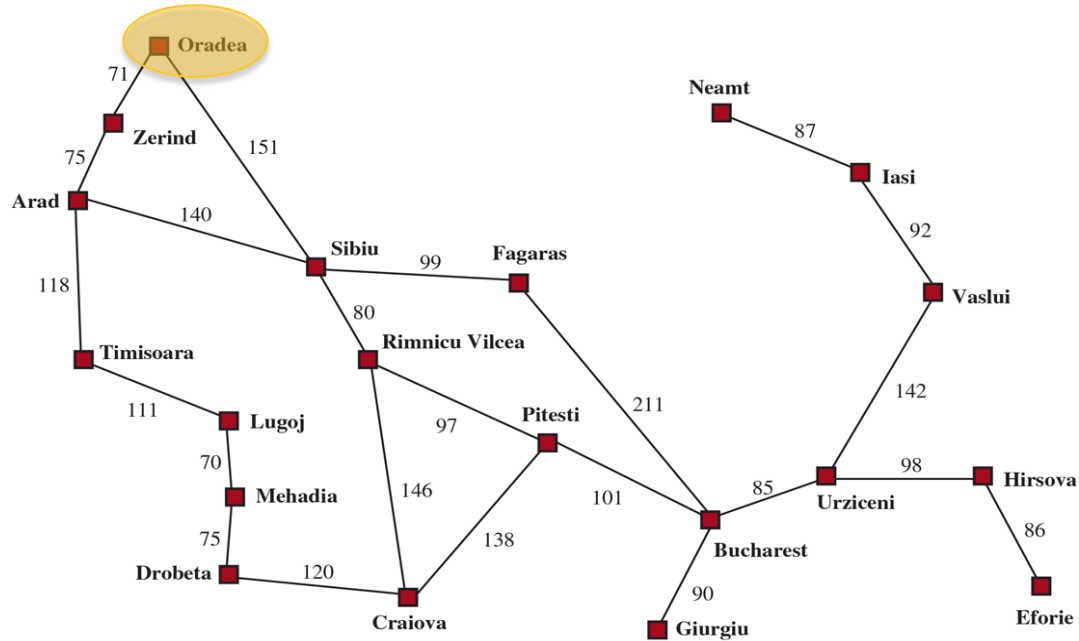


Fig 3.1, Russell & Norvig's Textbook

A State Space

Initial State

Goal State(s)

Actions

Transition Model

Action-Cost Function

Defining
the Search
Problem

$s = \text{Oradea},$
 $a = \text{ToZerind},$
 $s' = \text{Zerind}$
 $\text{cost}(s, a, s') = 71 \text{ miles}$

Example Problems



- Robo-Vacuum
- Sliding-Tile Puzzle
- Route Finding
- Automatic Motor Part Assembly



Formally Define the Search Problem for the Sliding-Tile Puzzle.



Search Algorithm

It is an algorithm that takes a search problem as an input and it returns a solution to the problem, or an indication that the problem has no solution.

Some search algorithms represent the problem as **a tree structure**. The root holds the **initial state**. The **nodes** represent different **states** that are reached by taking certain **actions**. Actions are represented by the **edges** connecting the nodes.

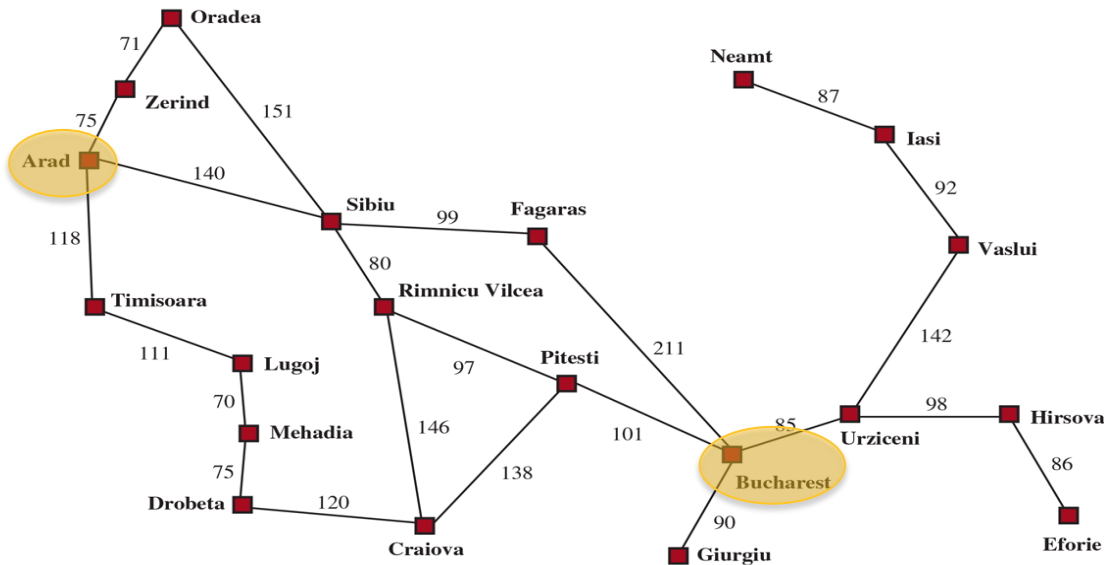


Fig 3.1, Russell & Norvig's Textbook

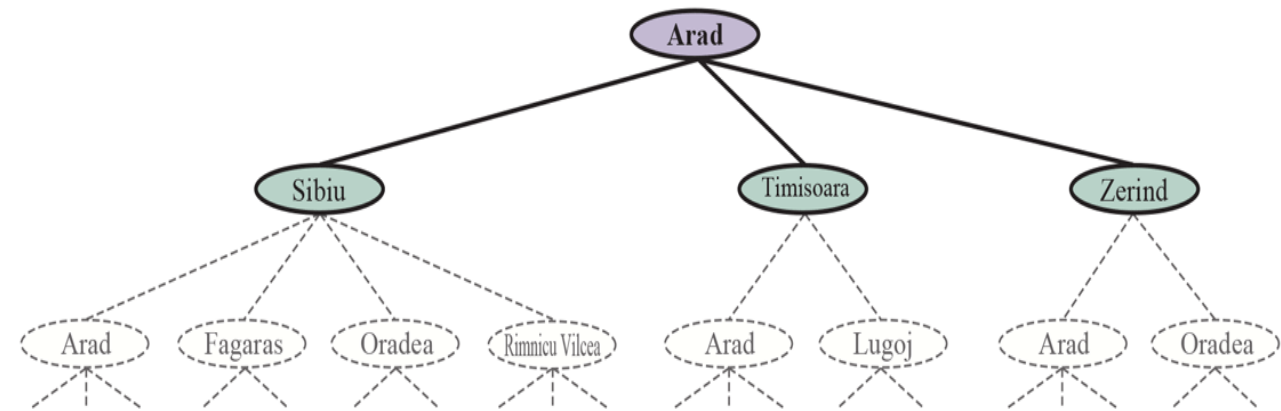


Fig 3.4, Russell & Norvig's Textbook

A search algorithm works by **expanding** nodes in the search tree. Each expanded node leads to a **new state**. The algorithm keeps **checking** if any of the expanded states is a **goal state**. The path from the initial state to a goal state is the **solution path**.





Fig 3.5, Russell & Norvig's Textbook



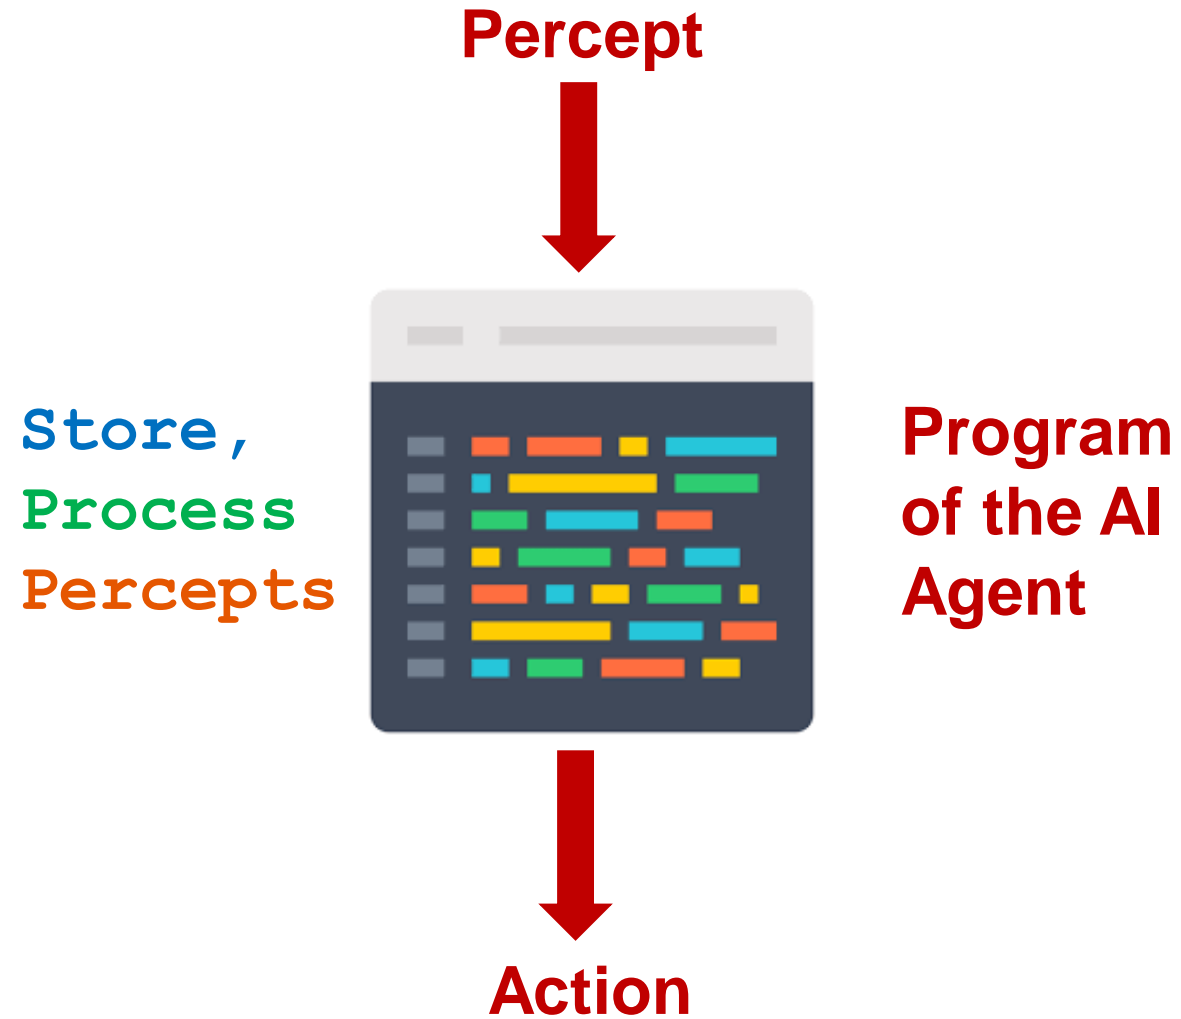
Which node should a search algorithm expand first??

Best-first Search



It is an algorithm that decides which node to expand based on some evaluation function $f(n)$. The node that has the minimum $f(n)$ is expanded first.

Best-first Search



Best-first Search

```
function BEST-FIRST-SEARCH(problem, f) returns a solution node or failure
  node  $\leftarrow$  NODE(STATE=problem.INITIAL)
  frontier  $\leftarrow$  a priority queue ordered by f, with node as an element
  reached  $\leftarrow$  a lookup table, with one entry with key problem.INITIAL and value node
  while not IS-EMPTY(frontier) do
    node  $\leftarrow$  POP(frontier)
    if problem.IS-GOAL(node.STATE) then return node
    for each child in EXPAND(problem, node) do
      s  $\leftarrow$  child.STATE
      if s is not in reached or child.PATH-COST < reached[s].PATH-COST then
        reached[s]  $\leftarrow$  child
        add child to frontier
  return failure

function EXPAND(problem, node) yields nodes
  s  $\leftarrow$  node.STATE
  for each action in problem.ACTIONS(s) do
    s'  $\leftarrow$  problem.RESULT(s, action)
    cost  $\leftarrow$  node.PATH-COST + problem.ACTION-COST(s, action, s')
    yield NODE(STATE=s', PARENT=node, ACTION=action, PATH-COST=cost)
```



**AI Agent
Program**

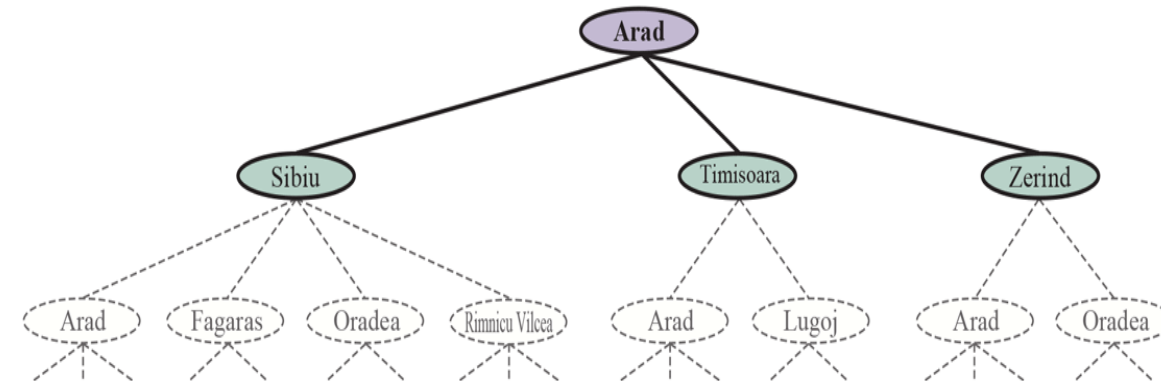


Fig 3.4, Russell & Norvig's Textbook

Fig 3.7, Russell & Norvig's Textbook

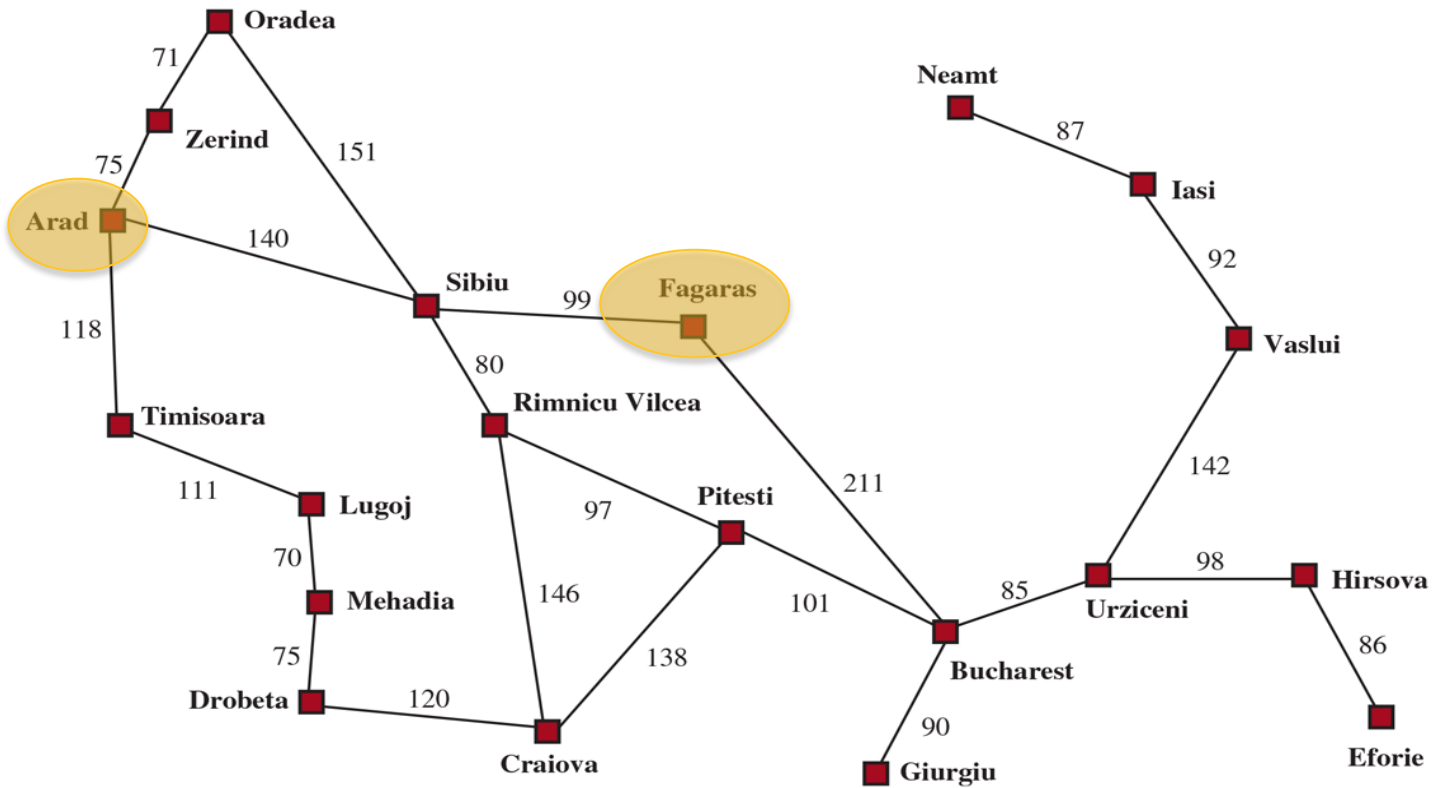


Fig 3.1, Russell & Norvig's Textbook

The search might return to state that was previously reached:

Arad → Sibiu → Arad

The search might reach one state through different paths:

Arad → Sibiu → Fagaras

Arad → Zerind → Oradea → Sibiu → Fagaras

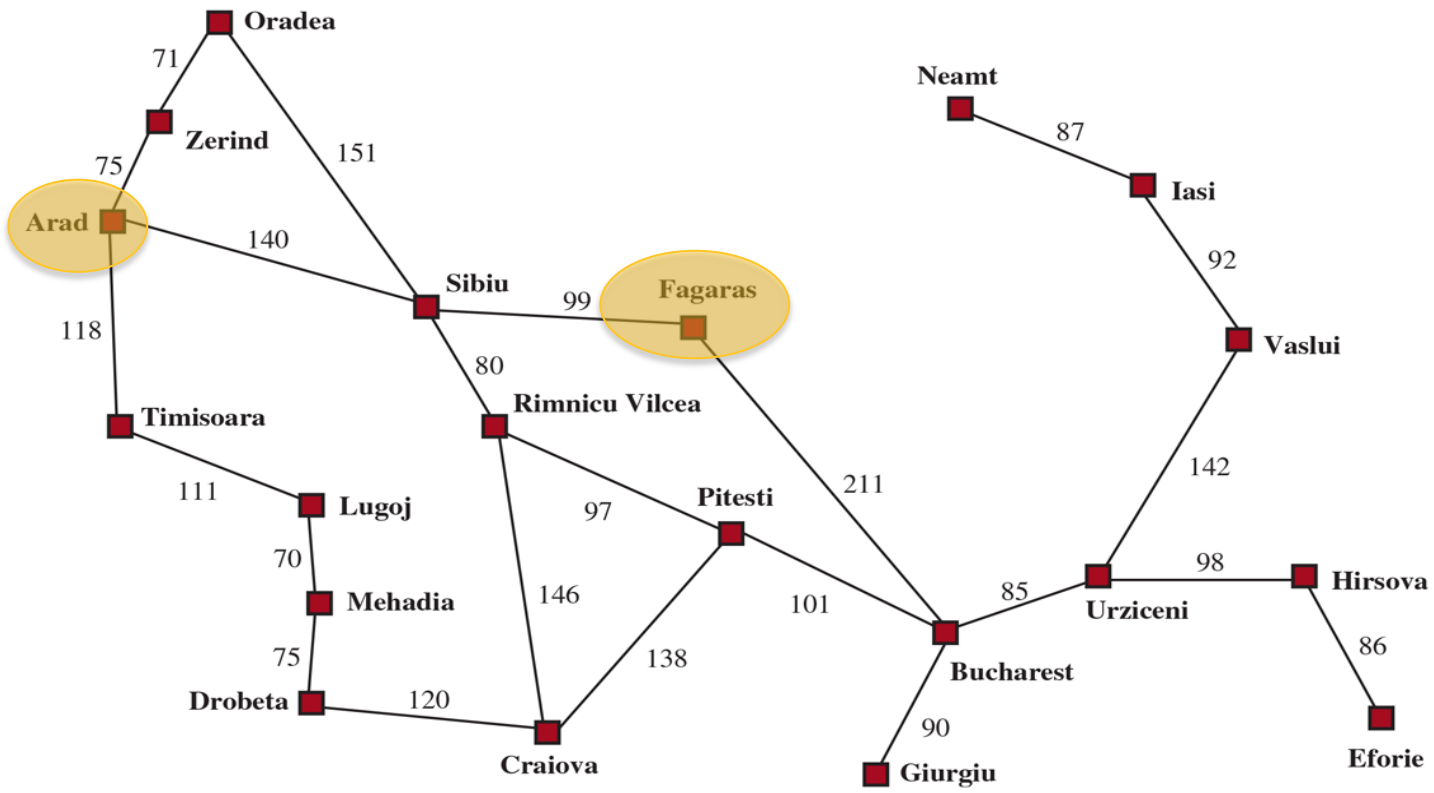


Fig 3.1, Russell & Norvig's Textbook

Arad → Zerind → Oradea

Cycles & Redundant Paths

Cycles & Redundant Paths

Cycles or redundant paths lead to infinite, useless search paths! The program may get into an infinite loop.

“Algorithms that can’t remember the past are doomed to repeat it.”

~ Ch3, Russell and Norvig’s

The algorithm should remember all visited states, detect redundant paths, keep only the most optimum path to each state.