

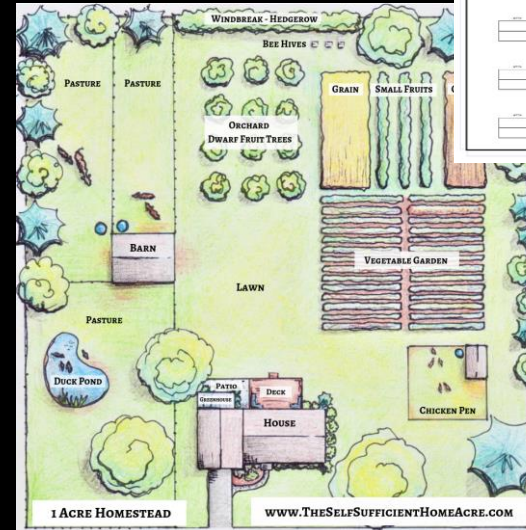
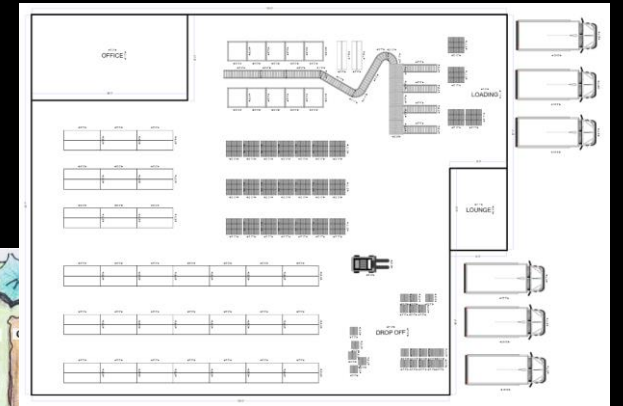

Search in Complex Environments

Asmaa Elbadrawy
PhD, Lecturer
IFT Program, ASU

**Given a problem,
find a solution state that
satisfies certain conditions
without worrying about the
path to the solution.**

Example Problems

- 8-Queen
- Integrated Circuit Design
- Factory Floor Layout
- Crop Planning



Local Search & Optimization Algorithms

- Hill Climbing
- Simulated Annealing

Local Search

A search from an initial state to neighboring states. It does not keep track of paths nor previously-reached states. It just keeps exploring neighboring states until a solution is found.

Optimization Problem

It is a problem in which we try to find the best state according to some objective function. We say it is required to maximize the objective value.

Optimization Problem




Fig 4.1, Russell & Norvig's Textbook

*The objective value changes from one state to the next. Maximizing the objective is done by moving up-hill until a peak is found. This is referred to as **Hill Climbing**.*

Hill Climbing starts from the initial state. With each step, it moves to the neighboring state with the highest value (called the direction of the **steepest ascent**). It terminates when a **peak** is reached. That is, there is no neighboring state that has a higher value. **It does not necessarily terminate at a global maximum.** It is not guaranteed to find the optimal solution (target state)!




Fig 4.1, Russell & Norvig's Textbook

Hill climbing gets stuck at a **local maximum** or a **plateau** 86% of the time!
Success rate=**14%**
If it is allowed to continue on a plateau, success rate increases to **94%**.




Fig 4.1, Russell & Norvig's Textbook

Another variation of Hill Climbing is **Random-Restarting Hill Climbing**. When the algorithm fails, it **randomly generates a new initial state** and **restarts**. On average, the algorithm restarts 7 times before it finds a solution.




Fig 4.1, Russell & Norvig's Textbook

Prohibiting downhill moves leads to getting stuck at a local maximum (non-solution state).

Random walking will eventually reach a solution but is very inefficient.

Combining both strategies can lead to a more efficient and complete algorithm.




Fig 4.1, Russell & Norvig's Textbook

Simulated Annealing

Annealing is a process used to temper or harden metals and glass by heating them to high temperatures then gradually cooling them down.

Simulated Annealing works with **cost**, not objective value. It tries to find the state with minimum cost. That is, it tries to find a **global minimum** point on the graph.

Simulated Annealing tries to “shake” the current state out of a local minimum to bounce it off towards, hopefully, a global minimum.

How is that done,
mathematically??
By introducing
controlled randomness




Fig 4.1, Russell & Norvig's Textbook



Temperature T decreases over time.

It specifies the probability of accepting a bad move.

Start at some random initial state

Repeat:

- Randomly generate a neighboring state
- If neighboring state is better than current state, accept
Else, accept with some probability p , that decreases with T